

CS 348

Intro to Artificial Intelligence

Day 5

Games

(slides based on Downy, Sood, Dan Klein, Pieter Abbeel)

Mike Rubenstein

Today:

- Class business
 - Lab 1 Due today, 7pm
- Adversarial search (chapter 5.1 - 5.4)
 - Alpha-beta pruning
 - Expecta-minimax
 - Evaluation functions
- Labs:
 - Intro Lab 3 (A*)
 - Discuss lab 2
 - Answer questions about lab 1

Minimax Example

```
def max-value(state):
```

```
    initialize  $v = -\infty$ 
```

```
    for each successor of state:
```

```
         $v = \max(v, \text{min-}$   
             $\text{value(successor)})$ 
```

```
    return v
```


```
def min-value(state):
```

```
    initialize  $v = +\infty$ 
```

```
    for each successor of state:
```


```
         $v = \min(v, \text{max-}$   
             $\text{value(successor)})$ 
```

```
    return v
```



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?



Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-beta pruning example

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
```

```
    initialize  $v = -\infty$ 
```

```
    for each successor of state:
```

```
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$ 
```

```
        if  $v \geq \beta$  return  $v$ 
```

```
         $\alpha = \max(\alpha, v)$ 
```

```
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):
```

```
    initialize  $v = +\infty$ 
```


```
    for each successor of state:
```

```
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$ 
```

```
        if  $v \leq \alpha$  return  $v$ 
```


```
         $\beta = \min(\beta, v)$ 
```

```
    return  $v$ 
```






Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
- With random ordering
 - Time complexity is $O(b^{3m/4})$
- This is a simple example of **metareasoning** (computing about what to compute)




Worst-Case vs. Average Case





Idea: Uncertain outcomes controlled by chance, not an adversary!

Expectimax Search

- Why wouldn't we know what the result of an action will be?
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the pacman ghosts respond randomly
 - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- **Expectimax search:** compute the average score under optimal play
 - Max nodes as in minimax search
 - Chance nodes are like min nodes but the outcome is uncertain
 - Calculate their **expected utilities**
 - I.e. take weighted average (expectation) of children



Expectimax Search



Expectimax Pseudocode

```
def value(state):
```

```
    if the state is a terminal state: return the state's utility
```

```
    if the next agent is MAX: return max-value(state)
```

```
    if the next agent is EXP: return exp-value(state)
```

```
def max-value(state):
```

```
    initialize v =  $-\infty$ 
```

```
    for each successor of state:
```

```
        v = max(v,  
               value(successor))
```

```
    return v
```

```
def exp-value(state):
```

```
    initialize v = 0
```

```
    for each successor of state:
```


```
        p = probability(successor)
```

```
        v += p * value(successor)
```

```
    return v
```


Expectimax Pseudocode

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```





$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

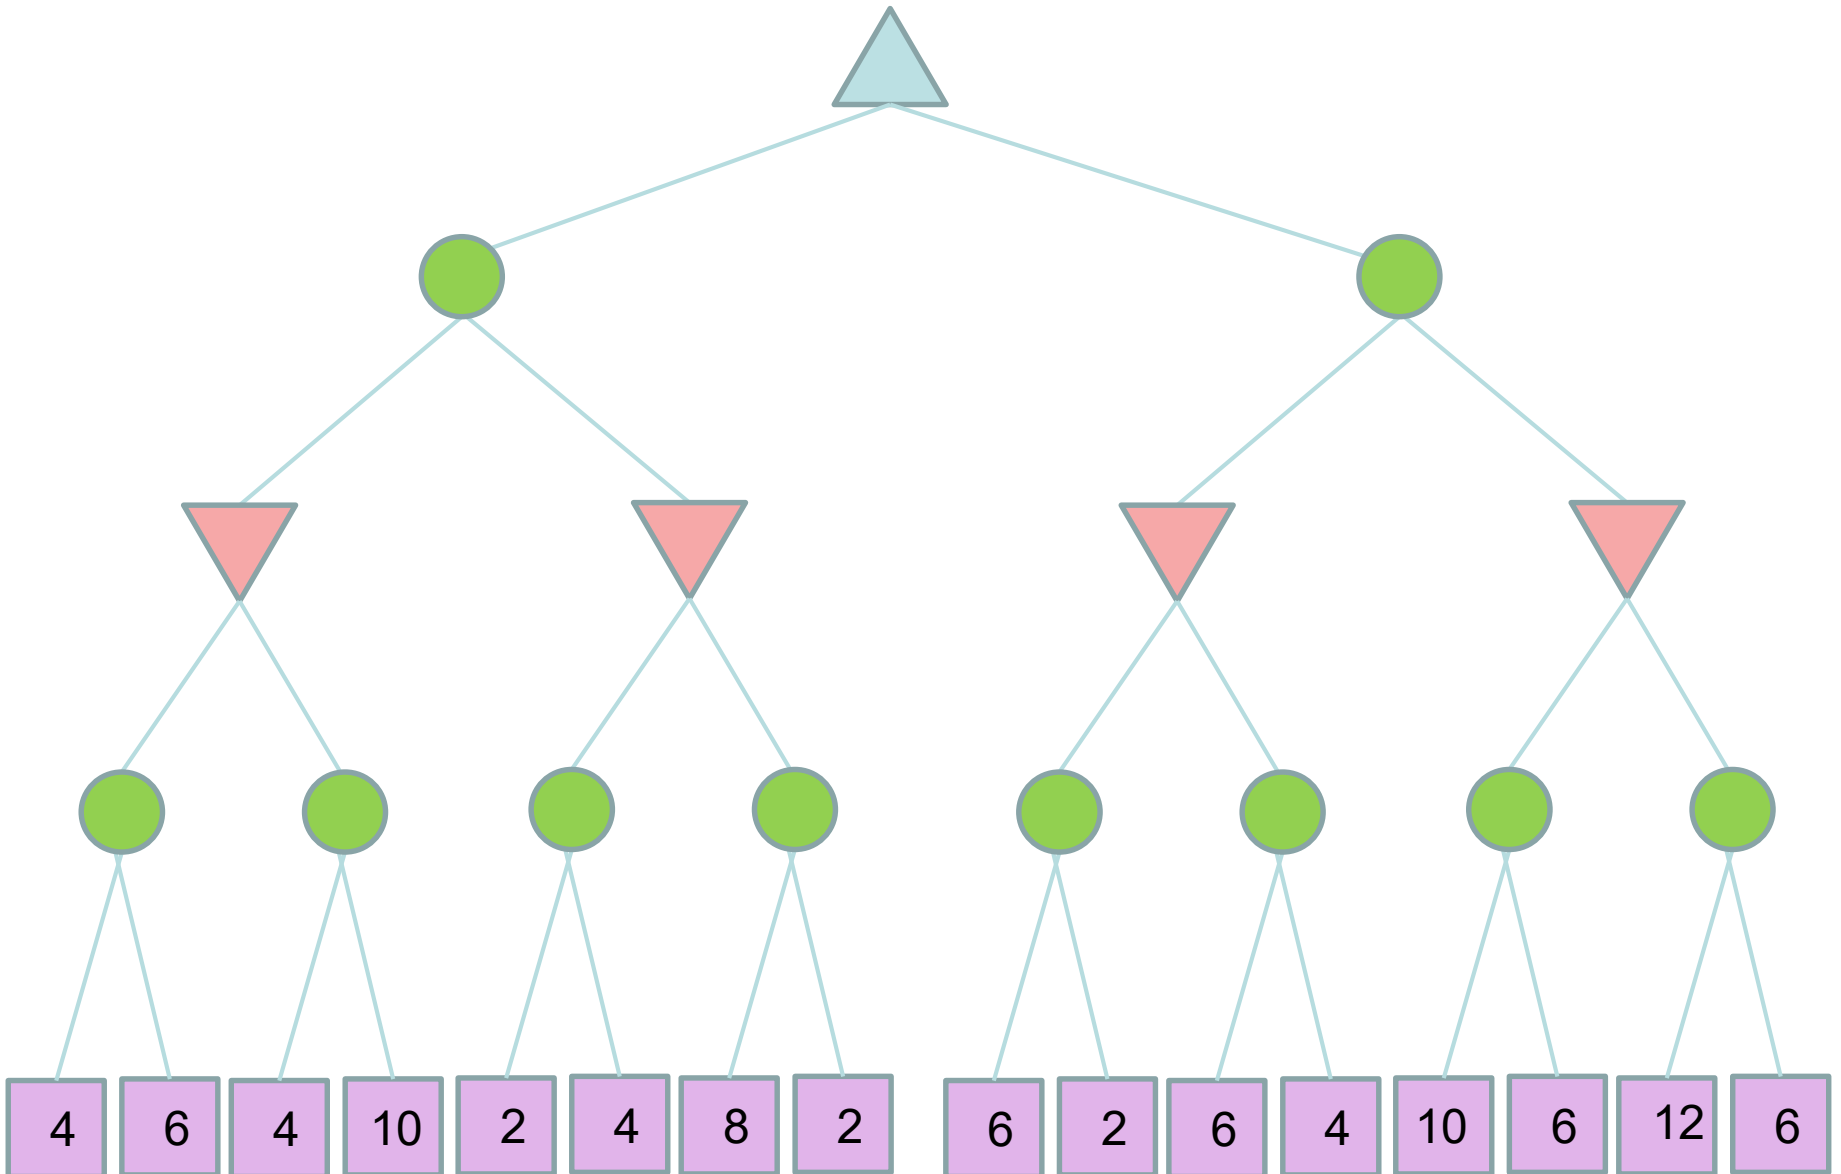
Expectimax Example



Expectimax Pruning?





Expecti-minimax





Utility values

- For worst-case minimax reasoning, terminal function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - We call this **insensitivity to monotonic transformations**



Utility values

- For worst-case minimax reasoning, terminal function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - We call this **insensitivity to monotonic transformations**
 - For average-case expectimax reasoning, we need *magnitudes* to be meaningful




Resource Limits





Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Instead, search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm




Depth Matters


- Evaluation functions tend to be imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation




Depth 2 search



Depth 10 search




Evaluation Functions



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search




- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.


Evaluation for Pacman



What features are we using?



Eval(n)=

Other Game Types



Mixed Layer Types

- E.g. Backgammon
- Expectiminimax
 - Environment is an extra “random agent” player that moves after each min/max agent
 - Each node computes the appropriate combination of its children






Example: Backgammon

- Dice rolls increase b : 21 possible rolls with 2 dice
 - Backgammon ≈ 20 legal moves
- As depth increases, probability of reaching a given search node shrinks
 - So usefulness of search is diminished
 - So limiting depth is less damaging
- Historic AI: TDGammon uses depth-2 search + very good evaluation function + reinforcement learning: world-champion level play
- 1st AI world champion in any game!



Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have utility tuples
 - Node values are also utility tuples
 - Each player maximizes its own component
 - Can give rise to cooperation and competition dynamically...



Summary

- Adversarial search
 - Mini-Max
 - Alpha-Beta pruning
 - Expecti-mini-max
 - Limited depth search

Lab 3 Intro


- A* search on grid world
 - Due April 28, 7pm
- Use Manhattan distance as $H(n)$
- In case of multiple nodes in frontier with lowest $F(n)$
 - Expand node with lowest x, if still tied, then expand node with lowest y
- No loops
- Exactly 1 starting position, 1 goal position
- OK to reuse portions of your code from lab 1

1	1	1	1	0
0	2	0	0	0
0	1	1	1	1
0	1	1	3	1
0	0	0	0	0

Lab 2, Due April 21, 7pm

■ 10-queens problem

- Queens constrained to column
- Find local minima using greedy search
 - Each step is motion of one queen in her column
 - 81 possible next moves
 - Follow tie breaking rules



Lab 2 discussion

Lab 2 discussion
